

# 3D N-Body Simulation with CUDA and ThreeJS

Mohammad A.Raji, Alok Hota

April 28, 2015

## 1 Introduction

N-body simulation, is a simulation of the effects of physical forces between particles in a dynamical system, and is used in many areas such as: astronomy, physics, chemistry, etc. The simulation is very time consuming due to the large number of particles and the need to calculate the force between all combinations.

In this work, we implement a serial and a parallel version of the simulation to show the effectiveness of GPU parallelization for the problem and also provide an interactive web interface for viewing the simulation results. Our simulation is tested with a maximum number of 1024 particles for 10000 timesteps.

## 2 Method

## 3 User Interface

To view the simulation results, we created a web interface that receives the simulation results as a **csv** file and renders it as a simulation in a 3D environment using ThreeJS - the WebGL based library.

The presented web interface provides the following the features:

- Panning and zooming in a 3D environment
- Pausing and playing the simulation
- Slow motion option
- Manual and automatic rotation option

A screenshot of our web interface is shown in Figure 1.

## 4 Tests

## 5 Evaluation

## 6 Conclusion

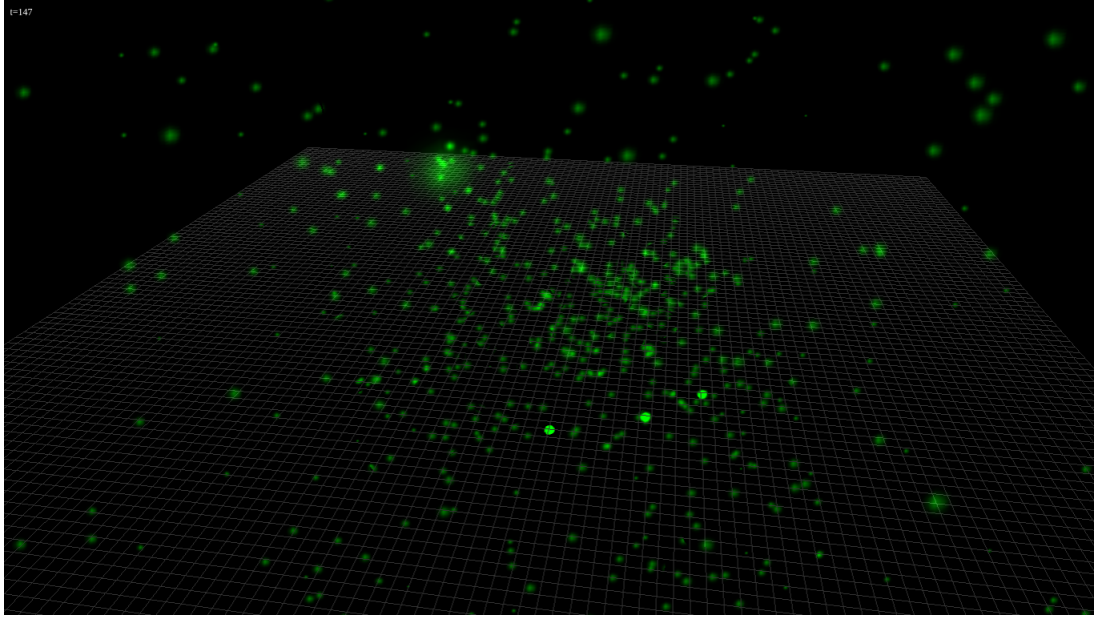


Figure 1: Web interface for our 3D n-body simulation

## 7 Appendix A: Header file

```
#include<stdio.h>
#include<stdlib.h>
#include<string.h>
#include<unistd.h>
#include<math.h>
#include<sys/time.h>
#include<time.h>

#define CMAX      1496           //+1 AU * 10e-5
#define CMIN      -1496
#define MMAX      9e2           //approximately mass of Ceres
#define MMIN      14e1         //approximately mass of Bennu
#define AMAX      1000
#define AMIN      -1000
#define VC        299792458     //speed of light

#define EPSILON2   0.5 f        //softener used to prevent r^2 -> 0

#define DEF_BODIES 256
#define G          1.0 f//6.673e-11f //gravitational constant
#define DEF_DELTA  0.1 f
#define DEF_STEPS  10000

//Lazy programming
int NUMBODIES, NUMSTEPS;
float DELTA_T;

//tools
int parse_args(int argc, char **argv);
int output = 1;

//parameter functions
float rand_acceleration();
```

```

float rand_coordinate();
float rand_mass();

//timing functions and struct
typedef struct {
    struct timeval start;
    struct timeval end;
} timer;
void start_timer(timer *t);
void stop_timer(timer *t);
float elapsed_time(timer *t);

```

## 8 Appendix B: Serial version

```

#include "nbody.h"

//CPU-specific structs and functions
typedef struct {
    float x;
    float y;
    float z;
} float3;

typedef struct {
    float x;
    float y;
    float z;
    float w;
} float4;
void interact(float4 *body_i, float4 *body_j, float3 *acc_i, float4
    ↪ *inter_i);

int main(int argc, char **argv) {
    //Get parameters, if any, from user
    NUMBODIES = DEF_BODIES;
    NUMSTEPS = DEF_STEPS;
    DELTA_T = DEF_DELTA;
    int status = parse_args(argc, argv);
    if(status)
        return status;

    int i, j, t;
    srand(time(NULL));
    timer perf_timer;
    timer total_timer;

    printf("Creating_bodies...\n");

    float4 *pos_mass;
    float4 *intermediate;
    float3 *acc;

    printf("Allocating_memory...\n");

    start_timer(&total_timer);

```

```

pos_mass      = (float4 *)malloc(NUMBODIES * sizeof(float4)
    ↪ );
intermediate = (float4 *)malloc(NUMBODIES * sizeof(float4)
    ↪ );
acc           = (float3 *)malloc(NUMBODIES * sizeof(float3)
    ↪ );

printf("Initializing bodies...\n");

for(i = 0; i < NUMBODIES; i++) {
    pos_mass[i].x = rand_coordinate();
    pos_mass[i].y = rand_coordinate();
    pos_mass[i].z = rand_coordinate();
    pos_mass[i].w = rand_mass();
}

printf("SIMULATION_SETTINGS:\n");
printf("__bodies__=%d\n", NUMBODIES);
printf("__steps__=%d\n", NUMSTEPS);
printf("__delta_t__=%f\n", DELTA_T);

printf("Running simulation...\n");

start_timer(&perf_timer);
for(t = 0; t < NUMSTEPS; t++) {
    for(i = 0; i < NUMBODIES; i++) {
        for(j = 0; j < NUMBODIES, j != i; j++) {
            interact(&pos_mass[i], &pos_mass[j]
                ↪ ), &acc[i], &intermediate[i])
                ↪ ;
        }
    }
    //Update positions
    for(i = 0; i < NUMBODIES; i++) {
        pos_mass[i] = intermediate[i];
    }
}
stop_timer(&perf_timer);
stop_timer(&total_timer);

printf("Simulation_runtime:\t%f_s\n", elapsed_time(&
    ↪ perf_timer));
printf("Total_runtime:\t%f_s\n", elapsed_time(&total_timer)
    ↪ );

if(output) {
    time_t raw_time;
    struct tm *current_time;
    time(&raw_time);
    current_time = localtime(&raw_time);
    char *filename = (char *)malloc(64);
    sprintf(filename, "cpu_%02d%02d%02d-%02d%02d%02d.
        ↪ nbd",
        current_time->tm_year%100, current_time->tm_mon,
        current_time->tm_mday, current_time->tm_hour,
        current_time->tm_min, current_time->tm_sec);

    printf("Saving to %s...\n", filename);
}

```



```

        printf("Usage: %s [-b num_bodies] [-s num_steps] [-t delta_t] [-o] \n",
               ↪ argv[0]);
        return 1;
    }
    else if(strcmp(argv[i], "-b") == 0) {
        NUMBODIES = atoi(argv[i + 1]);
    }
    else if(strcmp(argv[i], "-s") == 0) {
        NUMSTEPS = atoi(argv[i + 1]);
    }
    else if(strcmp(argv[i], "-t") == 0) {
        DELTA_T = atof(argv[i + 1]);
    }
    else if(strcmp(argv[i], "-o") == 0) {
        output = 0;
    }
    else {
        fprintf(stderr, "Error: unsupported flag %s\n", argv[i]);
        return -1;
    }
}
return 0;
}

float rand_coordinate() {
    return ((float)rand() / (float)RANDMAX) * (CMAX - CMIN) +
           ↪ CMIN;
}

float rand_acceleration() {
    return ((float)rand() / (float)RANDMAX) * (AMAX - AMIN) +
           ↪ AMIN;
}

float rand_mass() {
    return ((float)rand() / (float)RANDMAX) * (MMAX - MMIN) +
           ↪ MMIN;
}

void start_timer(timer *t) {
    gettimeofday( &(t->start), NULL);
}

void stop_timer(timer *t) {
    gettimeofday( &(t->end), NULL);
}

float elapsed_time(timer *t) {
    return (float) (t->end.tv_sec - t->start.tv_sec) +
           (t->end.tv_usec - t->start.tv_usec) /
           1000000.0;
}

```

## 9 Appendix C: Parallel version

```

#include "nbody.h"
#include <cuda.h>

//CUDA-specific vars and functions
__device__ int NBODIES;
__device__ float DT;
__global__ void main_nbody_kernel(float4 *dev_pos_mass, float3 *
    ↪ dev_acc,
float3 *dev_output, int cur_step);
__device__ void tile_nbody_kernel(float4 *my_pos_mass, float3 *
    ↪ my_acc);
__device__ void force_kernel(float4 *body_i, float4 *body_j,
float3 *acc_i);

int main(int argc, char **argv) {
    //Get parameters, if any, from user
    NUM_BODIES = DEF_BODIES;
    NUM_STEPS = DEF_STEPS;
    DELTA_T = DEF_DELTA;
    int status = parse_args(argc, argv);
    if(status)
        return status;
    cudaMemcpyToSymbol(NBODIES, &NUM_BODIES, sizeof(int), 0,
    cudaMemcpyHostToDevice);
    cudaMemcpyToSymbol(DT, &DELTA_T, sizeof(int), 0,
    cudaMemcpyHostToDevice);

    int i;
    srand(time(NULL));
    timer perf_timer;
    timer total_timer;

    printf("Creating_bodies...\n");

    float4 *host_pos_mass, *dev_pos_mass;
    float3 *host_acc, *dev_acc;
    float3 *host_output, *dev_output;

    printf("Allocating_host_memory...\n");

    start_timer(&total_timer);
    host_pos_mass = (float4 *)malloc(NUM_BODIES * sizeof(float4)
    ↪ );
    host_acc = (float3 *)malloc(NUM_BODIES * sizeof(float3)
    ↪ );
    host_output = (float3 *)malloc(NUM_BODIES * NUM_STEPS *
    ↪ sizeof(float3));

    printf("Allocating_device_memory...\n");

    cudaMalloc((void **)&dev_pos_mass, NUM_BODIES * sizeof(
    ↪ float4));
    cudaMalloc((void **)&dev_acc, NUM_BODIES * sizeof(float3));
    cudaMalloc((void **)&dev_output, NUM_BODIES * NUM_STEPS *
    ↪ sizeof(float3));

    printf("Initializing_bodies...\n");

```

```

for(i = 0; i < NUMBODIES; i++) {
    host_pos_mass[i].x = rand_coordinate();
    host_pos_mass[i].y = rand_coordinate();
    host_pos_mass[i].z = rand_coordinate();
    host_pos_mass[i].w = rand_mass();
}

printf("SIMULATION_SETTINGS:\n");
printf("  _bodies_ = %d\n", NUMBODIES);
printf("  _steps_ = %d\n", NUMSTEPS);
printf("  _delta_t_ = %f\n", DELTA_T);

/*
printf("Initial positions and masses:\n");
for(i = 0; i < NUMBODIES; i++) {
    printf("%d:\t%f\t%f\t%f\n", i, host_pos_mass[i].x,
        ↪ host_pos_mass[i].y,
        host_pos_mass[i].z, host_pos_mass[i].w);
}
*/

printf("Copying_to_device...\n");

cudaMemcpy(dev_pos_mass, host_pos_mass, NUMBODIES * sizeof
    ↪ (float3),
cudaMemcpyHostToDevice);
cudaMemcpy(dev_acc, host_acc, NUMBODIES * sizeof(float3),
cudaMemcpyHostToDevice);
cudaMemcpy(dev_output, host_output, NUMBODIES * NUMSTEPS
    ↪ * sizeof(float3),
cudaMemcpyHostToDevice);

printf("Running_kernel...\n");

start_timer(&perf_timer);
int block_size = (NUMBODIES < 16) ? 4 : (NUMBODIES < 256)
    ↪ ? 16 : 32;
int grid_size = NUMBODIES / block_size;
int mem_size = (block_size+1) * sizeof(float4);
printf("KERNEL_SETTINGS:\n");
printf("  _bodies_ = %d\n", NUMBODIES);
printf("  _tile_size_ = %d\n", block_size);
printf("  _grid_size_ = %d\n", grid_size);
for(i = 0; i < NUMSTEPS; i++) {
    main_nbody_kernel<<<grid_size, block_size, mem_size
        ↪ >>>(dev_pos_mass,
        dev_acc, dev_output, i);
}
stop_timer(&perf_timer);

printf("Simulation_runtime:\t%f_s\n", elapsed_time(&
    ↪ perf_timer));

printf("Copying_to_host...\n");

cudaMemcpy(host_output, dev_output, NUMBODIES * NUMSTEPS
    ↪ * sizeof(float3),

```





```

//iterate over all tiles and update position and
    ↪ acceleration
//each iteration loads one tile's worth of data from global
    ↪ memory
//these reads should be coalesced
int i, tile;
for(i = 0, tile = 0; i < NBODIES; i += blockDim.x, tile++)
    ↪ {
        //index into global for this thread's body *for
            ↪ this tile*
        int tile_id = tile * blockDim.x + threadIdx.x;

        //threads collaborate to load from global for this
            ↪ tile
        tile_pos_mass[threadIdx.x] = dev_pos_mass[tile_id];
        __syncthreads();

        //update acceleration for this thread's body for
            ↪ this tile
        tile_nbody_kernel(&my_pos_mass, &my_acc);
        __syncthreads();
    }

    //update position for this body
    my_pos_mass.x += my_acc.x;
    my_pos_mass.y += my_acc.y;
    my_pos_mass.z += my_acc.z;

    //update global position array
    dev_pos_mass[global_id] = my_pos_mass;
    dev_acc[global_id] = my_acc;

    //update global output
    dev_output[cur_step * NBODIES + global_id].x = my_pos_mass.
        ↪ x;
    dev_output[cur_step * NBODIES + global_id].y = my_pos_mass.
        ↪ y;
    dev_output[cur_step * NBODIES + global_id].z = my_pos_mass.
        ↪ z;
}

__device__ void tile_nbody_kernel(float4 *my_pos_mass, float3 *
    ↪ my_acc) {
    //tile position array from the outer kernel
    //pre-loaded with this tile's positions and masses
    extern __shared__ float4 tile_pos_mass[];

    //iterate over each body in the tile and calculate its
        ↪ effect on
    //this thread's body
    int i;
    for(i = 0; i < blockDim.x; i++) {
        force_kernel(my_pos_mass, &tile_pos_mass[i], my_acc
            ↪ );
    }
}

```

```

__device__ void force_kernel(float4 *body_i, float4 *body_j, float3
    ↪ *acc_i) {
    //calculate distance components
    float3 d;
    d.x = body_j->x - body_i->x;
    d.y = body_j->y - body_i->y;
    d.z = body_j->z - body_i->z;

    //use epsilon softener
    // r^2 + epsilon^2
    float denominator = d.x * d.x + d.y * d.y + d.z * d.z +
    ↪ EPSILON2;
    //cube and sqrt to get (r^2 + epsilon^2)^(3/2)
    denominator = sqrt( denominator * denominator * denominator
    ↪ );

    float acc = G * body_j->w / denominator;

    //update acceleration
    acc_i->x += acc * d.x * DT;
    acc_i->y += acc * d.y * DT;
    acc_i->z += acc * d.z * DT;
}

int parse_args(int argc, char **argv) {
    int i;
    for(i = 1; i + 1 <= argc; i += 2) {
        if(strcmp(argv[i], "-h") == 0) {
            printf("Usage: %s [-b num_bodies] [-s
    ↪ num_steps] [-t delta_t] [-o] \n", argv[0]);
            return 1;
        }
        else if(strcmp(argv[i], "-b") == 0) {
            NUMBODIES = atoi(argv[i + 1]);
        }
        else if(strcmp(argv[i], "-s") == 0) {
            NUMSTEPS = atoi(argv[i + 1]);
        }
        else if(strcmp(argv[i], "-t") == 0) {
            DELTA_T = atof(argv[i + 1]);
        }
        else if(strcmp(argv[i], "-o") == 0) {
            output = 0;
        }
        else {
            fprintf(stderr, "Error: unsupported flag %s
    ↪ \n", argv[i]);
            return -1;
        }
    }
    return 0;
}

float rand_coordinate() {
    return ((float)rand() / (float)RANDMAX) * (CMAX - CMIN) +
    ↪ CMIN;
}

```

```

float rand_acceleration() {
    return ((float)rand() / (float)RAND_MAX) * (AMAX - AMIN) +
        ↪ AMIN;
}

float rand_mass() {
    return ((float)rand() / (float)RAND_MAX) * (MMAX - MMIN) +
        ↪ MMIN;
}

void start_timer(timer *t) {
    gettimeofday( &(t->start), NULL);
}

void stop_timer(timer *t) {
    gettimeofday( &(t->end), NULL);
}

float elapsed_time(timer *t) {
    return (float) (t->end.tv_sec - t->start.tv_sec) +
        (t->end.tv_usec - t->start.tv_usec) /
        1000000.0;
}

```